

What I found in the data.

I am building SASA Intelligence Lab on the SASA datasets in the Data Lake. On 2 September I shared a list of everything that got in my way. This is the same list, checked again against the live API on **11 September**. For each item it shows what I reported then and what I see now. It also covers what is new since then, and how I have handled these problems in the app so far.

Where things stand on 11 September

Two items are fixed. Gobardhan can be read again, and every dataset that was missing is now live. One item is mostly fixed and four are partly fixed. **Seventeen are not fixed yet.** One can no longer be tested, because the filter it depended on has been removed.

Nothing in the data has been revised since my previous pull on 10 September. The catalogue has the same 42 datasets, every small table is identical row for row, the CDMA daily files still cover 12 to 27 August only, and PR door to door still stops at 7 August.

The most serious new problem is in how the data comes out of the API. **Neither paging nor the export route returns the CDMA daily files complete.** Both return the right number of rows, but about half of those rows are repeats. That is N1 below.

I also found mistakes in my own 2 September list. They are corrected at the end, with what went wrong.

FIXED 2 B3 B8	MOSTLY FIXED 1 B2	PARTLY FIXED 4 A2 A8 B7 C1
--	--	---

NOT FIXED 17 A1 A3 A4 A5 A6 A7 A9 A10 A11 A12 A13 A14 B4 B5 B6 C2 C3	CANNOT TEST 1 B1	NEW 7 N1 N2 N3 N4 N5 N6 N7
---	---	---

Checked: all 42 datasets in the catalogue, live, 11 September 2026

Previous version: 2 September 2026, 33 datasets

Owners: ● Department ● TCS ● NextGen ● Either, needs routing

AT A GLANCE

Every item and its status

I have put a suggested owner on each item. I do not know your internal split well enough to be sure, so where it could sit with more than one team it says so.

ID	ISSUE	SUGGESTED OWNER	11 SEPTEMBER
NEW SINCE 2 SEPTEMBER			
N1	Paging and export do not return the CDMA daily files complete	● NextGen or TCS	NEW
N2	Monthly history removed, and an old key now returns a different table	● TCS	NEW
N3	One table served under several dataset names	● TCS	NEW
N4	Some new LGD codes put a town in the wrong district	● TCS	NEW
N5	Markapuram has no district code in the CDMA daily files	● TCS	NEW
N6	Place names spilling into a count column	● Dept or TCS	NEW
N7	PR door to door stops at 7 August and stores every day three times	● Dept or TCS	NEW
A • THE DATA			
A1	The same town has different names in different datasets	● Dept or TCS	NOT FIXED
A2	<code>district_id</code> is not reliable	● TCS	PARTLY FIXED
A3	<code>ulb_code</code> holds the mandal code	● TCS	NOT FIXED

ID	ISSUE	SUGGESTED OWNER	11 SEPTEMBER
A4	No way to tell a real zero from not reported	● Dept	NOT FIXED
A5	Month number and month name do not match	● Dept or TCS	NOT FIXED
A6	No way to tell monthly from cumulative	● Dept	NOT FIXED
A7	Same field, different names, in files shared together	● TCS	NOT FIXED
A8	No record ID, submission date or revision number	● TCS	PARTLY FIXED
A9	Units and period formats are not stated	● Dept	NOT FIXED
A10	Fields with no documentation	● TCS	NOT FIXED
A11	Datasets do not say what they should cover	● Dept or TCS	NOT FIXED
A12	Documentation mistakes, now in the live data	● TCS	NOT FIXED
A13	Most daily collection values are zero and none are blank	● Dept	NOT FIXED
A14	The same place and month twice, with different numbers	● Dept or TCS	NOT FIXED
B • THE API			
B1	A month filter and a whole read disagree	● TCS	CANNOT TEST
B2	Wrong filter names in the docs, and errors that say nothing	● TCS	MOSTLY FIXED
B3	Gobardhan could not be read	● NextGen or TCS	FIXED
B4	Page size stuck at 100, token lasts 300 seconds	● NextGen or TCS	NOT FIXED
B5	Allowed purposes are not published	● NextGen or TCS	NOT FIXED
B6	No way to check whether a dataset changed	● NextGen or TCS	NOT FIXED
B7	The same filter has different names	● TCS	PARTLY FIXED
B8	Documented datasets not yet live	● NextGen	FIXED
C • DECISIONS			
C1	Agree on LGD codes, and which LGD code	● All three	PARTLY FIXED
C2	Agree what a zero means	● Dept	NOT FIXED
C3	Check the data before it is accepted	● Dept or TCS	NOT FIXED

NEW SINCE 2 SEPTEMBER

Problems that were not on the first list

I found these while checking the old list. Most of them came in with the platform revision in the first week of September.

N1 Paging and export do not return the CDMA daily files complete

● NextGen or TCS **NEW**

The row count is always right. The rows are not. Both ways of getting the data out return the correct number of rows, but about half of them are repeats standing in for rows that never arrive. Nothing in the response says so.

`msw_door_to_door_collection_api`, checked 11 September.

The file should hold one row per secretariat per day. $4,023 \text{ secretariats} \times 16 \text{ days} = 64,368$, which is exactly the total the API reports.

Export, one call per district for all 26 district codes: 64,368 rows back, only **32,206** different. Some rows arrive up to 8 times.

Paging through the whole file on 8 September: 64,368 rows back, **40,636** different.

The two routes miss different rows. In Visakhapatnam export gave 4,614 different rows and paging gave 5,067, and only 2,741 were in both.

Paging again does not help. A second full pass over Visakhapatnam found no new rows at all. I exported six districts a second time. **The same rows came back, repeated the same number of times, in all six.** This is not random. The export gives the same half every time: 1,439 of 2,880 rows in Anantapuram, 4,614 of 9,216 in Visakhapatnam.

Waste segregation and wet waste processing lose rows the same way. In the same six districts, 12,507 of 25,024 and 12,539 of 25,088 rows were different, 50.0% each.

Everything I could get by every route put together: **50,634 of 64,368 rows (78.7%)**.

The gaps are not spread evenly. 1 row is missing for 13 August and 2,858 are missing for 27 August.

`identification_of_bulk_waste_generators_api` behaves the same way: 64,528 rows back by export, 32,286 different, 50,777 of 64,528 (78.7%) by every route together.

What would fix this: every row served exactly once, by paging and by export. My guess is that the query underneath has no fixed sort order, so each page is cut from a different ordering, but I cannot confirm that from outside. Until it is fixed, any total built from these files is built on a partial file.

N2 Monthly history was removed, and an old key now returns a different table

● TCS **NEW**

ODF Plus model villages: 84 rows for May, June and July in my 28 August copy. Now 28 rows and no month columns at all.

Plastic waste management units: the same change, 84 rows down to 28 with no months.

E-autos service model: 166 rows down to 84.

FSTPs: the old key `sasa_sac_establishing_fstps_information_api` used to return 70 ULB rows of plant capacity. It now returns 28 district rows of targets and achievement. The plant data moved to `fstps_stps_cotreatment_new1_api`.

The old keys still answer, so anyone reading them gets different data under the same name, with no notice.

What would fix this: keep the monthly history, or publish it separately. If a key starts returning a different table, give it a new key.

N3

One table served under several dataset names

● TCS **NEW**

The door to door e-autos, push carts and tri-cycles keys, and the machinery e-autos service model key, all return the **same 84 rows**. The API names `d2d_collection_vehicles_api` as the table behind all four.

The three 50 percent green programme keys return the **same 400 rows**, from `apgbc_info_new1_api`. `soak_pits_api` returns all **336 Compost Pits rows**, identical, and its own `work_name` column says Compost Pits.

`waste_egregation_api` is the door to door collection file without the `collected_households` column. In the six districts I exported, all 25,024 segregation rows match a collection row column for column.

Read as separate programmes, these get counted two, three or four times.

What would fix this: one key per table. Where keys have to share a table, say so in the catalogue.

N4

Some of the new LGD codes put a town in the wrong district

● TCS **NEW**

The new LGD columns look as if they were matched on name, so a town can land on a mandal with the same name in another district.

Guduru, listed under Kurnool, is coded as Guduru mandal in **Krishna** (`swacch_survekshan_info_new1_api`).

Gudur, listed under SPSR Nellore, is coded as Gudur mandal in **Kurnool** (same dataset).

Ramachandrapuram, listed under Konaseema, is coded as Ramachandrapuram mandal in **Tirupati**, in both the Swachh Survekshan table and the 50 percent green programme table.

Rajampet carries two different mandal codes, 1120 and 5246, inside the same green programme table.

Some other changes, such as Giddalur to Markapuram and Kandukur to Prakasam, match the district reorganisation, so I have not counted those as errors.

What would fix this: check the name matching against LGD before it goes into every dataset. A wrong code is harder to catch than a missing one.

N5

Markapuram has no district code in the CDMA daily files

● TCS **NEW**

In both daily files `district_code` is the text **NULL** for Markapuram: 816 rows each, covering Markapur, Kanigiri, Giddalur and Podili. Every other district has a number. Filtering on

district_code: "NULL" does return them.

Markapuram became a district on 31 December 2025, so it looks as if the new district was never given a code in this feed.

What would fix this: give Markapuram its LGD district code in the CDMA feeds.

N6

Place names spilling into a count column

● Dept or TCS **NEW**

In `identification_of_bulk_waste_generators_api`, `no_of_bwgs` should be a number. For some Vijayawada secretariats it holds text:

```
secretariat_name: "62-Fire Station Area" , no_of_bwgs: " Prakash Nagar" , i_ts: "1"
secretariat_name: "02-Machavaram Masi" , no_of_bwgs: "id Area" , i_ts: "0"
```

The secretariat name has been split at a comma or cut short, and the rest has been pushed into the next columns. In the rows I could retrieve there are 18 counts like this and 15 that are missing.

What would fix this: quote text fields properly when the file is loaded. The same fix would also remove the stray quote marks in A1.

N7

PR door to door stops at 7 August and stores every day three times

● Dept or TCS **NEW**

July: $13,351 \text{ gram panchayats} \times 31 \text{ days} \times 3 \text{ copies} = \mathbf{1,241,643}$, exactly the total the API reports.

August: $13,351 \times 7 \text{ days} \times 3 \text{ copies} = \mathbf{280,371}$, again exactly the reported total.

On 11 September August still holds only 1 to 7 August, more than five weeks after the month ended. There is no September data. May, June and July are unchanged.

What would fix this: one copy of each panchayat day, and the rest of August. If the three copies are deliberate, say what makes them different.

SECTION A • THE DATA

Issues in the data itself

What is in the files, and how it is named and coded. Some of this can be fixed in processing. The parts about what was never recorded cannot, and those are marked as the department's.

A1

The same town has different names in different datasets

● Dept or TCS **NOT FIXED**

2 SEPTEMBER

ULB names changed from file to file. Atmakur had 10 spellings and Gudur had 10. Two names came back with quote marks inside the value, "JANGAREDDY GUEDEM" and "YSR TADIGADAPA".

11 SEPTEMBER

Every example is still there, and the new datasets added more. **Atmakur now has 14 spellings and Gudur 13.** Quote marks now appear inside **17 names across 10 datasets**, for example "Palasa Kasibugga" and "GVMC VISAKHAPATNAM". There is still no ULB ID.

What would fix this: a permanent ULB ID on every dataset, old and new. The counts are higher partly because the catalogue grew from 33 to 42 datasets.

A2

district_id is not reliable

● TCS

PARTLY FIXED

2 SEPTEMBER

The same `district_id` pointed to different districts. id 17 was Bapatla and Kakinada.

11 SEPTEMBER

Inside each dataset it is now consistent, except id 20, which is both NTR and VMC in the three MEPMA datasets. Across datasets it still means different places. LGD district codes have been added beside it, which is the better key.

id 17 is **Anakapalli** in compost pits, magic drains and soak pits, **Bapatla** in community sanitary complexes and SBM IHHL, **Prakasam** in the CDMA e-waste and plastic ban datasets, and **Kakinada** in the three MEPMA datasets.

What would fix this: use the LGD district code everywhere and drop `district_id`, or make it mean one thing.

A3

ulb_code holds the mandal code, not a ULB code

● TCS

NOT FIXED

2 SEPTEMBER

487 of 500 rows checked had `ulb_code` equal to `api_lgd_mandal_code`. GVMC showed that a mandal code cannot identify a ULB that covers several mandals.

11 SEPTEMBER

Checked on every row I could retrieve: **31,362 of 32,206 (97.4%)** still match the mandal code. GVMC still has `ulb_code` 1086 alongside four mandal codes: 1086, 1087, 1089 and 4865. No dataset has an LGD urban local body code.

What would fix this: the LGD urban local body code. This is going into every dataset now, so it is cheaper to correct before that than after.

A4

There is no way to tell a real zero from not reported

● Dept **NOT FIXED**

2 SEPTEMBER

Missing values showed up four ways: `0`, empty text, `null`, or no field at all. In the CDMA daily data not one value was blank, so a non report looked exactly like a zero.

11 SEPTEMBER

Unchanged. No dataset has a code for not reported, and no collection value in the CDMA daily files is blank or null. PR door to door does mark each day `IS_COLLECTED: Yes` or `No`, but it has no code for a day nobody reported either.

What would fix this: one code for not reported, different from zero, used the same way everywhere. See C2.

A5

Month number and month name do not match

● Dept or TCS **NOT FIXED**

2 SEPTEMBER

All 35 July FSTP rows had month number 7 and month name JUNE. SERP had `"July "` with a trailing space. Case varied between datasets.

11 SEPTEMBER

The same 35 rows are unchanged, and they have been copied into the replacement dataset `fstps_stps_cotreatment_new1_api`. Each SERP dataset still has 3 rows with `"July "`. Case still varies: May and may, June, JUNE and june, July, july and JULY.

What would fix this: one standard period field, and values in a row that agree with each other.

A6

There is no way to tell if a number is monthly or cumulative

● Dept **NOT FIXED**

2 SEPTEMBER

SERP circular economy showed 11,084 against a target of 5,000, which is 221%, because the field is cumulative while similar datasets are monthly.

11 SEPTEMBER

The same 11,084 against 5,000. One of the three SERP catalogue descriptions now says cumulative. The other two do not, and nothing in the data marks it.

What would fix this: say for each measure whether it is monthly or cumulative, and give the denominator for every percentage.

A7

The same field has different names in files shared together

● TCS **NOT FIXED**

2 SEPTEMBER

Door to door collection used `sachivalayam_code` . Bulk waste generators and wet waste processing used `secretariat_code` for the same values.

11 SEPTEMBER

Unchanged, and waste segregation arrived using the `sachivalayam` spelling. Two daily files use one name and two use the other.

What would fix this: one agreed list of field names before a release.

A8

No record ID, submission date or revision number

● TCS

PARTLY FIXED

2 SEPTEMBER

0 of 32 readable datasets had a record ID, a submission date or a revision number.

11 SEPTEMBER

Four datasets that arrived later carry insert and update times (`i_ts` , `u_ts`): community sanitary complexes, housing IHHL, SBM IHHL and sewage treatment. Sewage treatment also has `rec_id` . **No dataset has a revision number**, and none of the older datasets gained any of these.

The formats already differ. Sewage treatment gives "14-08-2026 14:09" with quote marks inside the value. Housing IHHL gives 2026-08-28 16:14:52 .

What would fix this: a record ID, submission date and revision number on every row, in one format. This is also what would settle A14.

A9

Units and period formats are not stated

● Dept

NOT FIXED

2 SEPTEMBER

From the documentation: `garbage_segregation` had no unit, a month of `202701` looked like a future date, and `achivement` was misspelled.

11 SEPTEMBER

Now visible in the live data. `garbage_segregation` still has no unit, though its values (300, 650, 1000) look like household counts. Compost pits, magic drains and soak pits run from `202604` to `202703` , one row per district per month of the financial year, with a target every month and achievement only up to `202608` . That explains `202701` , but it is written down nowhere. `achivement` is still misspelled.

What would fix this: a unit and a denominator for each measure, and a line explaining the period format.

A10

Fields with no documentation

● TCS NOT FIXED

2 SEPTEMBER

`active_indicator` and `i_ts` were undocumented. I withdrew `a_in`, `s_no` and `u_ts` because no live dataset had them.

11 SEPTEMBER

`active_indicator` is still there and is 1 on every row. `i_ts` is still undocumented. **The three fields I withdrew now exist** in the four datasets that arrived later. The integration guide I have does not describe any of them.

What would fix this: document these fields or take them out.

A11

Datasets do not say what they should cover

● Dept or TCS NOT FIXED

2 SEPTEMBER

If a ULB is missing from a dataset there is no way to know whether it has nothing to report or never submitted.

11 SEPTEMBER

Unchanged. The catalogue has no field for expected coverage.

What would fix this: each dataset states what it covers, for example all 123 ULBs, so a gap is a fact and not a guess.

A12

Documentation mistakes, now in the live data

● TCS NOT FIXED

2 SEPTEMBER

The waste segregation key was spelled two ways, and the Soak Pits sample in the guide was a copy of Compost Pits.

11 SEPTEMBER

`waste_segregation_api` spelled correctly returns **404 dataset_not_found**. The live key is the misspelled `waste_egregation_api`. The Soak Pits copy is no longer only in the guide: the live `soak_pits_api` returns the Compost Pits rows (N3).

What would fix this: the correct key, and real Soak Pits data.

A13

Most daily collection values are zero and none are blank

● Dept NOT FIXED

2 SEPTEMBER

From 23 samples of 100 rows across the file: 93% of `collected_households` values were zero and none were blank.

11 SEPTEMBER

The closest I can get to a full count is **95.2% zero: 48,180 of the 50,634** different secretariat days I could retrieve. It cannot be a complete count, because the API does not return the whole file (N1). Still no blank or null values, and still only 12 to 27 August.

What cannot be told apart: whether collection is not happening, which would be serious, or whether it is not being reported. The app does show urban collection for one reference day, 12 August, and counts a reported zero as no collection. I label it as one reported day and not a service rating, and I will not treat it as performance until this is answered. My 2 September wording about districts was also wrong, see the corrections at the end.

A14

The same place and month comes back twice, with different numbers

● Dept or TCS

NOT FIXED

2 SEPTEMBER

ODF Plus had every district twice for May, and 10 pairs disagreed. Alluri Sitharama Raju was 1,296 or 389. Sweeping machines had Nellore as 19 and 4 machines in June and again in July.

11 SEPTEMBER

ODF Plus has no duplicates now, but only because the table was replaced by 28 rows with no months and different targets (N2). Alluri Sitharama Raju's target went from 1,323 to 2,115. The May question was never answered. **Sweeping machines is unchanged:** Nellore is still 19 and 4 in both months.

Identical copies still in the data: legacy waste 3 pairs, IHHL 2 pairs, and the new sewage treatment dataset 6 pairs. PR SWPC operators lists every one of its 28 districts twice.

What would fix this: someone who knows the source says which Nellore row is right. After that, A8 would stop this happening again.

SECTION B • THE API

Issues with the platform

These are about getting data out, not about the data. N1 belongs here too and is the most serious of them.

B1

A month filter and a whole read disagree

● TCS

CANNOT TEST

2 SEPTEMBER

In IHHL, six ULBs came back as June in a whole read and as July when filtered to July.

11 SEPTEMBER

The month filter has been removed. The new IHHL dataset accepts only `lgd_dist_code`, and the old key now serves the same new table, so this test cannot be run any more. I cannot tell whether the problem underneath was fixed.

B2

Wrong filter names in the docs, and errors that say nothing

● TCS **MOSTLY FIXED**

2 SEPTEMBER

The catalogue listed filters for legacy waste that returned 502. A wrong filter name gave a 502 with no hint of the right one.

11 SEPTEMBER

The catalogue now lists `month_id` and `year`, which work. A wrong filter name now gets a **400 that names the allowed filters**. A valid filter with a value that matches nothing, for example `month_id: "99"`, still gets a bare 502 instead of an empty result.

```
400: "Unknown filter field(s) for tableKey
'sasa_100_percent_clearance_of_legacy_waste_api': __probe__. Allowed: month_id, year."
502: "The Data Lake API could not complete this request."
```

What would still help: an empty result, not a 502, when nothing matches.

B3

Gobardhan could not be read

● NextGen or TCS **FIXED**

2 SEPTEMBER

Every route returned 502, including export, while 32 other datasets answered.

11 SEPTEMBER

It returns **28 district rows**. Only 4 of the 28 districts report any achievement. The response names the table `gobardhanunits_api`, not the key I asked for.

B4

Page size stuck at 100, token lasts 300 seconds

● NextGen or TCS **NOT FIXED**

2 SEPTEMBER

`pageSize` was ignored without a warning, and a large pull outlived its token.

11 SEPTEMBER

Asking for `pageSize` 2 or 1,000 still returns 100. The access token still lasts 300 seconds, though a refresh token now lets a long pull carry on. There is an export route that returns a whole filtered result in one call, which helps, but it loses rows the same way paging does (N1).

What would fix this: a page size that can be set, and an export that returns each row once.

B5

Allowed purposes are not published

● NextGen or TCS

NOT FIXED

2 SEPTEMBER

Each dataset accepts different purposes and the catalogue did not say which.

11 SEPTEMBER

The catalogue still does not list them. The error message does: "Purpose 'AUDIT' is not allowed for tableKey 'compost_pits_api'. Allowed: BENEFIT_DISBURSEMENT."

What would fix this: the allowed purposes for each dataset in `GET /api/v1/datasets`.

B6

No way to check whether a dataset changed

● NextGen or TCS

NOT FIXED

2 SEPTEMBER

The only way to know was to download everything again.

11 SEPTEMBER

Responses now carry an ETag header, but **it changes on every call**, even when the rows are identical, so it cannot show a change. Six datasets now say which older dataset they replace, which helps with keys but not with content.

What would fix this: a content hash for each dataset and month, or an ETag that only changes when the rows do.

B7

The same filter has different names

● TCS

PARTLY FIXED

2 SEPTEMBER

District had three filter names on 8 of 33 datasets. Month had four. Only `year` was consistent, on 27.

11 SEPTEMBER

District filtering now works on 32 of 42 datasets, but under six names. Month still has four. `year` is offered on only 10 datasets, and one of those spells it `YEAR`.

District: `lgd_dist_code` 9, `lgd_district_code` 8, `district_name` 5, `district_id` 4, `district_code` 4, `DISTRICT_ID` 2
Month: `month` 5, `month_id` 4, `month_number` 3, `MONTH_ID` 1

What would fix this: one name per concept, for example `district`, `month`, `year`.

B8

Documented datasets not yet live

● NextGen **FIXED**

2 SEPTEMBER

Five CDMA datasets and three PR datasets were documented but not on the portal.

11 SEPTEMBER

All eight are live. The catalogue went from 33 to 42 datasets. The PR keys carry a date, for example `sasa_pr_no_of_swpcs_operationalised_api_27_aug_2026`, and waste segregation is live under the misspelled key (A12).

SECTION C • DECISIONS

Things that need to be agreed together

These are conventions rather than one team's job. Once they are settled, a lot of the rest follows.

C1

Agree on LGD codes, and on which LGD code

● All three **PARTLY FIXED**

2 SEPTEMBER

LGD was partly in the data under two conventions, and there was no LGD urban local body code.

11 SEPTEMBER

LGD district and mandal codes are now in many datasets. There is still no LGD urban local body code. The district code appears under three column names, and the old name columns were kept, so one row can carry `dstrt_nm`, `district_name` and `lgd_district_name`. Some of the new codes are wrong (N4).

LGD district code columns: lgd_dist_code, lgd_district_code, api_lgd_dist_code

Still to agree: the ULB code, one column name for each code, and whether the old name columns go.

C2 Agree what a zero means

● Dept **NOT FIXED**

2 SEPTEMBER

Without this decision no collection percentage can be worked out from the daily data.

11 SEPTEMBER

No decision yet, and the daily files still have no way to mark a missing report (A13).

Decision needed: one code for not reported, different from zero, used everywhere.

C3 Check the data before it is accepted

● Dept or TCS **NOT FIXED**

2 SEPTEMBER

A schema check at submission would catch much of Section A.

11 SEPTEMBER

The Soak Pits copy, the FSTP month mismatch, quote marks in names, text in a count column and a district coded as NULL all reached the live API.

Suggestion: validate at ingestion and send the errors back to whoever submitted.

WHAT I DID WITH THE DATA

How I built SASA Intelligence Lab around these problems

I could not wait for every item above to be fixed, so the app handles them itself. None of this changes a source value. It decides what gets counted, what gets held out, and what the app refuses to say. The item references show which problem each rule answers.

KEEPING THE DATA

I keep my own copy of every dataset

B6 A8

The app reads saved copies, not the live API. Each copy records the table key, response ID, filters, retrieval time, row counts and a fingerprint of its content, so I can see when a dataset changes and accept each new version on purpose. On 10 September I kept 36 of the 42 datasets in the catalogue: 27 unchanged and 9 revised.

I kept the older copies that still have history

N2

ODF Plus, plastic waste units and FSTPs lost their months in the September revision. My 28 August copies still have May to July, so I kept them beside the new versions instead of replacing them.

A pull is complete only when the rows are

N1 N7

Getting as many rows as the API reports proves only that enough requests were made. I count different records, not rows returned, and I state coverage as what was retrieved against what should exist.

COUNTING

Identical duplicates are counted once

A14 N7

Where the same place and period appears more than once with the same values, it counts once. For PR door to door that took **3,965,247** rows down to **1,285,814** panchayat days across May to August.

Conflicting duplicates are held out of every total

A14 A8

Where the same place and period has two different values, both rows are left out of totals and both numbers are shown, because nothing in the data says which is right. Nellore's 19 and 4 sweeping machines are shown that way.

Datasets that serve the same table count once

N3 A12

I detect this from the rows, not the names. Soak Pits is excluded from totals because it returns Compost Pits rows, and the three green programmes each read their own column from the shared table. If the platform separates them again, the flag drops off on the next pull.

Missing is never zero

A4 A13 A11

Every measure keeps separate counts for reported, zero and missing. A rate is not shown unless its denominator was reported, and a place with no record shows as **Not returned**, never as 0.

Months are compared on the same days

N7

August stops at day 7, so I compare May to August on days 1 to 7 only, for a fixed group of **85,769** panchayat days that have a valid value in every month. On that basis reported rural collection went from **74.71% to 73.42%**. I do not compare months of different lengths on their own totals.

Sundays and second Saturdays are shown separately

A13

Sundays report 1.8% to 3.0% statewide and second Saturdays about 11%, against 82% to 94% on working days. Those are scheduled non collection days, not failures. The app shows the rate with every day counted and with working days only (**86.82% to 85.24%**), says which one is on screen, and lists the days it left out.

Capacity is not output

A9

TPD and KLD figures are shown as configured capacity, never as waste processed. Legacy waste cleared is shown as reported, not as verified.

IDENTITY AND COMPARISON

A name match is a candidate, not an identity

A1 A2 A3

Matching two spellings of a town is a guess, so the app never treats it as an approved join. Anything that depends on it stays unscored. Today no ULB profile is scored.

I use the LGD codes as the source's claim, and report where they disagree

N4 C1

Where a row states its own LGD code, I can use that as evidence. Where the codes contradict each other, the app lists the disagreement for review. It never fills in a code for a row that does not have one.

Two sources are combined only if they agree

A6 A11

When I joined two sources on district and month, **84 district months** lined up and **none** agreed on the target. The app shows that result and refuses to add the two together.

The one cross source reading that passes its checks

A7 A3

Door to door collection and waste segregation both key on the numeric secretariat code, and in my 8 September pull they report the same 3,024 codes for 12 August, with the same household counts. That opens three of the five checks the app requires. The other two stay closed: one day is not a basis for judging performance, and no scoring policy has been approved. That view shows positions, not a rank. Gap Radar does rank ULBs for the same day, as described below.

Data I hold back

A5 A13

The 35 FSTP rows with conflicting months are kept but not scored. The 2024 Swachh Survekshan results are kept apart from 2026 operations. Neither is treated as performance.

Urban daily collection is shown for one day

A13 N1

Gap Radar places and ranks **103 of 118 ULBs** on collection reach and segregation for one reference day, 12 August, against draft 80% lines that are not approved targets. The rest are held out because their codes need review, or they have no households or no collection reported. A reported zero counts as no collection, which is exactly the problem in A13. That day came from my 8 September pull, which I now know held **3,024 of the 4,023** secretariats for 12 August (N1). Today's export has all 4,023, and collection for the day moves from **56.9% to 56.1%**. The app still shows the 8 September figures.

Six you can reproduce in a few minutes

Each of these can be run in the Data Lake playground. I ran every one on 11 September.

N1 **The export returns repeats instead of rows.** Export `msw_door_to_door_collection_api` with the filter `district_code: "NULL"`.

It returns 816 rows. Count the different pairs of `date1` and `sachivalayam_code`: there are 409. The file should have 816, one for each secretariat on each of 16 days. Run the same export again and the same 409 come back, each repeated the same number of times.

N3 **Soak Pits is Compost Pits.** Query `soak_pits_api` and `compost_pits_api` with no filter.

Both return 336 rows. The rows are identical, and every Soak Pits row says `work_name: "Compost Pits"`.

A14 **Two answers for Nellore.** Query `sasa_sac_sweeping_machines_information_api` with no filter.

Nellore appears four times: 19 machines and 4 machines for JUNE, and 19 and 4 again for JULY. Nothing on the rows says which is current.

A5 **July rows labelled June.** Query `fsteps_stps_cotreatment_new1_api` with no filter.

All 35 rows with `mnth_no "7"` have `mnth_nm "JUNE"`.

B2 **A clear error and an unclear one.** Query `sasa_100_percent_clearance_of_legacy_waste_api` twice.

With filter `{ "__probe__": "x" }`: 400, naming the allowed filters `month_id` and `year`.
With filter `{ "month_id": "99", "year": "2026" }`: 502 `upstream_request_failed`.

A1 **Names with quote marks, and Bobbili.** Query `sasa_mepma_entrepreneurs_promoted_for_circular_economy_api` with no filter.

`ulb_name` includes "JANGAREDDY GUDEM" and "YSR TADIGADAPA" with the quote marks inside the value, and BOBBILLI, which other datasets spell Bobbili.

ACTIONS

What would help

Everything still open, grouped by who is most likely to own it, most urgent first. Fixed items are left out.

● If only three things get picked up

in this order

1. **Make paging and export return every row once.** Every figure built from the daily files depends on it. N1
2. **Decide what a zero means.** Until then no collection rate can be published from the daily data. C2 A13
3. **Fix the codes before they spread.** The LGD urban local body code in place of the mandal code, and a check on the name matched LGD codes. A3 N4

● Department

about what was recorded

1. **One code for not reported,** different from zero. A4 C2
2. **The rest of August for PR door to door,** and September. N7
3. **Monthly or cumulative** for each measure, with its denominator. A6
4. **Units, and the period format** for compost pits, magic drains and soak pits. A9

● Department or TCS

depends where the field is set

1. **Say which Nellore sweeping machine row is right,** and whether the ODF Plus May history can come back. A14 N2
2. **A permanent ULB ID** on every dataset. A1
3. **One period field,** with values that agree inside a row. A5
4. **What each dataset should cover.** A11
5. **Validation and proper quoting** at submission. C3 N6

● TCS

processing, naming and the API contract

1. **A code for Markapuram** in the CDMA feeds. N5
2. **Keep monthly history**, and a new key when a key changes table. N2
3. **One key per table**, or say which keys share one. N3
4. **Record ID and revision number** on every row. A8
5. **Drop or fix district_id**. A2
6. **One filter name per concept**. B7
7. **An empty result, not a 502**, when a filter value matches nothing. B2
8. **One field name per concept**. A7
9. **Document or remove** active_indicator, i_ts, a_in, s_no and u_ts. A10
10. **The correct waste segregation key**, and real Soak Pits data. A12

● NextGen or TCS

platform behaviour

1. **Every row once**, by paging and by export. N1
2. **A page size that can be set**. B4
3. **A change signal** that only changes when the rows do. B6
4. **Allowed purposes** in the catalogue. B5

● What I will do

so this is not only a list of requests

- Check every item again after each change and write down what is fixed, as I did here.
- Share the exact request and response for anything on this list.
- Label any figure from the daily files as partial until N1 is fixed.
- Keep my saved copies, so history removed from the API is not lost.

CORRECTIONS

What I got wrong

I would rather list my own mistakes than leave them for someone else to find. These are the ones in my 2 September list.

WRONG ON 2 SEPTEMBER

Bobbilli does exist

My 2 September corrections said Bobbilli does not exist in the live data. **That was wrong**. BOBBILLI is in the three MEPMA datasets, both in my 28 August copy and live today. My check was case sensitive and missed the all capitals spelling. The original A1 example was right.

WRONG ON 2 SEPTEMBER

No district is zero on every row

I wrote that eleven of twenty three sample points were zero on every row, and the actions listed Konaseema, Visakhapatnam, West Godavari and Nandyal as zero districts. The sample points were slices of 100 rows, not whole districts. Counted over every row I could retrieve, no district is entirely zero.

OUT OF DATE

The PDF and the link did not match

The PDF I shared on 2 September still said 26 districts in A2. The link had already been corrected to 28. This version replaces both.

CHANGED SINCE 2 SEPTEMBER

Some counts are higher because there is more data

The catalogue grew from 33 to 42 datasets, so counts such as spellings per town are higher partly for that reason, not only because things got worse. I have said so where it matters.

CARRIED OVER

Earlier corrections that still stand

Rows were not being moved between months over time. The record count on filtered requests was not wrong. The state has had 28 districts since 31 December 2025.

SASA Intelligence Lab · Data issues

Checked live 11 September 2026 · first shared 2 September 2026